

Problem Set 4: Biased Methods and Regularization

Donggyun Kim
27008257

3/21/2018

1) Properties of PLS Regression

a)

$$||\mathbf{w}_h|| = \mathbf{w}_h^t \mathbf{w}_h = 1.$$

$$\mathbf{p}_h = \frac{\mathbf{X}_{h-1}^t \mathbf{z}_h}{\mathbf{z}_h^t \mathbf{z}_h} = \frac{\mathbf{X}_{h-1}^t \mathbf{X}_{h-1} \mathbf{w}_h}{\mathbf{w}_h^t \mathbf{X}_{h-1}^t \mathbf{X}_{h-1} \mathbf{w}_h}.$$

$$\mathbf{w}_h^t \mathbf{p}_h = \frac{\mathbf{w}_h^t \mathbf{X}_{h-1}^t \mathbf{X}_{h-1} \mathbf{w}_h}{\mathbf{w}_h^t \mathbf{X}_{h-1}^t \mathbf{X}_{h-1} \mathbf{w}_h} = 1.$$

b)

$$\mathbf{X}_h = \mathbf{X}_{h-1} - \mathbf{z}_h \mathbf{p}_h^t = \mathbf{X}_{h-1} - \mathbf{X}_{h-1} \mathbf{w}_h \mathbf{p}_h^t.$$

$$\mathbf{X}_h \mathbf{w}_h = \mathbf{X}_{h-1} \mathbf{w}_h - \mathbf{X}_{h-1} \mathbf{w}_h \mathbf{p}_h^t \mathbf{w}_h.$$

$$\mathbf{w}_h^t \mathbf{X}_h^t = (\mathbf{X}_h \mathbf{w}_h)^t = \mathbf{w}_h^t \mathbf{X}_{h-1}^t - \mathbf{w}_h^t \mathbf{p}_h \mathbf{w}_h^t \mathbf{X}_{h-1}^t = \mathbf{w}_h^t \mathbf{X}_{h-1}^t - \mathbf{w}_h^t \mathbf{X}_{h-1}^t = 0.$$

2) Bias of Regression Coefficients in PCR

a)

$$\hat{\beta}_Z^{(k)} = V_k \hat{\beta}_{PCR}^{(k)}.$$

$$\mathbb{E}[V_k \hat{\beta}_{PCR}^{(k)}] = V_k (Z_k^t Z_k)^{-1} Z_k^t \mathbb{E}[y] = V_k (\Lambda_k)^{-1} (X V_k)^t (X \beta) = V_k (\Lambda_k)^{-1} V_k^t X^t X \beta.$$

$$\mathbb{E}[V_k \hat{\beta}_{PCR}^{(k)} - \beta] = \mathbb{E}[V_k \hat{\beta}_{PCR}^{(k)}] - \beta = V_k (\Lambda_k)^{-1} V_k^t X^t X \beta - \beta = (V_k (\Lambda_k)^{-1} V_k^t X^t X - I) \beta.$$

b)

$$\hat{\beta}_Z^{(p)} = V_p \hat{\beta}_{PCR}^{(p)}.$$

$$\mathbb{E}[V_p \hat{\beta}_{PCR}^{(p)}] = V_p (Z_p^t Z_p)^{-1} Z_p^t \mathbb{E}[y] = V_p ((X V_p)^t X V_p)^{-1} (X V_p)^t (X \beta) = V_p (V_p^t X^t X V_p)^{-1} V_p^t X^t X \beta = V_p V_p^t (X^t X)^{-1} V_p V_p^t X^t X \beta = \beta.$$

$$\mathbb{E}[V_p \hat{\beta}_{PCR}^{(p)} - \beta] = 0.$$

3) Bias of Ridge Regression Coefficients

$$X = UDV^t.$$

$$\hat{\beta}_r = (X^t X + kI)^{-1} X^t y = ((UDV^t)^t UDV^t + kI)^{-1} (UDV^t)^t y = (VD^2 V^t + kI)^{-1} VDU^t y.$$

$$\mathbb{E}[\hat{\beta}_r] = (VD^2 V^t + kI)^{-1} VDU^t X \beta = (VD^2 V^t + kI)^{-1} VDU^t UDV^t \beta = (VD^2 V^t + kI)^{-1} VD^2 V^t \beta.$$

$$D^2 = \Lambda.$$

$$\mathbb{E}[\hat{\beta}_r - \beta] = \mathbb{E}[\hat{\beta}_r] - \beta = (V\Lambda V^t + kI)^{-1} V\Lambda V^t \beta - \beta = ((V\Lambda V^t + kI)^{-1} V\Lambda V^t - I)\beta.$$

4) Models for Solubility Data

4.1) PCR

```
library(AppliedPredictiveModeling)
library(caret)
library(pls)
library(elasticnet)
library(ggplot2)

data(solubility)

# 10-fold cross-validation
ctrl <- trainControl(method = "cv", number = 10)

set.seed(1991)
pcr_fit <- train(x = solTrainXtrans, y = solTrainY,
                 method = "pcr",
                 tuneLength = 40,
                 trControl = ctrl,
                 preProcess = c("center", "scale"))


pcr_fit

## Principal Component Analysis
##
## 951 samples
## 228 predictors
##
## Pre-processing: centered (228), scaled (228)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 855, 857, 855, 857, 857, 855, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      2.0393844  0.01330618  1.5768251
##   2      1.9802436  0.08618857  1.5600094
##   3      1.7098785  0.30775962  1.3508311
##   4      1.6039381  0.38902366  1.2491550
##   5      1.5655620  0.41438454  1.2160916
##   6      1.4419644  0.50548830  1.1133483
##   7      1.2816012  0.60844627  1.0008166
```

```

##      8      1.2837318  0.60644396  1.0030821
##      9      1.2845568  0.60575993  1.0043077
##     10      1.2630788  0.61833048  0.9805825
##     11      1.2389026  0.63310862  0.9595090
##     12      1.2411704  0.63155212  0.9651889
##     13      1.2394323  0.63270101  0.9608776
##     14      1.1923383  0.66024430  0.9291284
##     15      1.1673741  0.67410421  0.9143805
##     16      1.1192399  0.69994038  0.8829070
##     17      1.0568173  0.73286887  0.8342530
##     18      1.0469036  0.73821332  0.8264794
##     19      1.0369036  0.74282876  0.8135459
##     20      1.0106942  0.75501073  0.7931272
##     21      1.0084359  0.75693801  0.7910714
##     22      1.0077410  0.75762006  0.7907431
##     23      0.9824906  0.76956926  0.7709785
##     24      0.9823339  0.76968098  0.7703055
##     25      0.9807215  0.77067359  0.7701617
##     26      0.9744918  0.77358933  0.7643884
##     27      0.9646638  0.77820852  0.7572924
##     28      0.9635618  0.77853798  0.7549437
##     29      0.9622526  0.77899192  0.7528061
##     30      0.9626567  0.77883793  0.7531227
##     31      0.9464459  0.78605250  0.7422126
##     32      0.9308610  0.79410333  0.7278662
##     33      0.9205904  0.79876222  0.7171017
##     34      0.9210315  0.79901759  0.7171882
##     35      0.9118633  0.80247076  0.7125940
##     36      0.8955050  0.80969558  0.6990453
##     37      0.8898150  0.81229196  0.6925830
##     38      0.8869713  0.81345299  0.6896258
##     39      0.8880278  0.81262668  0.6904784
##     40      0.8815142  0.81589128  0.6832985
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 40.
# plot the RMSEs against the number of PCs
plot(pcr_fit)

```



the number of PCs that gives the minimum RMSE value

```
pcr_fit$bestTune
```

```
##      ncomp
```

```
## 40      40
```

Make a plot of the regression coefficient paths

```
pcr_coef <- pcr_fit$finalModel$coefficients
```

```
n <- nrow(pcr_coef)
```

```
p <- as.numeric(pcr_fit$bestTune)
```

```
variable <- rep(rownames(pcr_coef), p)
```

```
PC <- rep(1:p, each = n)
```

```
pcr_coef <- as.matrix(pcr_coef)
```

```
dat <- data.frame(variable,
                  coefficient = pcr_coef,
                  PC,
                  stringsAsFactors = FALSE)
```

```
ggplot(dat, aes(x = PC, y = coefficient, col = variable)) +
  geom_step() +
  xlab("# Component") +
  ylab("Coefficient") +
  theme(legend.position = "none")
```



4.2) PLSR

```
set.seed(1991)
pls_fit <- train(x = solTrainXtrans, y = solTrainY,
  method = "pls",
  tuneLength = 30,
  trControl = ctrl,
  preProcess = c("center", "scale"))


pls_fit
```

```
## Partial Least Squares
##
## 951 samples
## 228 predictors
##
## Pre-processing: centered (228), scaled (228)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 855, 857, 855, 857, 857, 855, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      1.2834799  0.6039945  0.9916002
##   2      1.0565748  0.7313375  0.8338149
##   3      0.9421131  0.7896398  0.7295970
```

```

##      4      0.8672854  0.8202404  0.6763828
##      5      0.8234981  0.8385276  0.6400779
##      6      0.7858546  0.8533009  0.6099090
##      7      0.7576326  0.8632534  0.5803275
##      8      0.7462298  0.8670690  0.5718116
##      9      0.7336407  0.8715244  0.5620758
##     10      0.7197659  0.8770720  0.5536012
##     11      0.7168387  0.8779665  0.5497677
##     12      0.7125968  0.8796722  0.5472087
##     13      0.7077955  0.8812824  0.5413748
##     14      0.7072088  0.8816105  0.5404051
##     15      0.7020926  0.8833477  0.5355129
##     16      0.7027772  0.8829757  0.5367203
##     17      0.6980264  0.8849057  0.5346358
##     18      0.7001560  0.8843897  0.5370392
##     19      0.6987062  0.8852104  0.5342642
##     20      0.6999291  0.8847274  0.5341114
##     21      0.6995291  0.8851151  0.5357248
##     22      0.7009689  0.8846909  0.5347051
##     23      0.7038482  0.8835812  0.5356718
##     24      0.7060065  0.8828946  0.5373671
##     25      0.7103322  0.8814441  0.5406134
##     26      0.7140135  0.8801809  0.5413654
##     27      0.7170543  0.8792074  0.5427062
##     28      0.7175550  0.8790947  0.5423010
##     29      0.7205393  0.8781144  0.5432047
##     30      0.7202488  0.8781004  0.5422940
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 17.
# plot the RMSEs against the number of PLS components
plot(pls_fit)

```



the number of PLS components that gives the minimum RMSE value

```
pls_fit$bestTune
```

```
##      ncomp
```

```
## 17      17
```

make a plot of the regression coefficient paths

```
pls_coef <- pls_fit$finalModel$coefficients
```

```
n <- nrow(pls_coef)
```

```
p <- as.numeric(pls_fit$bestTune)
```


```
variable <- rep(rownames(pls_coef), p)
```

```
comp <- rep(1:p, each = n)
```

```
pls_coef <- as.matrix(pls_coef)
```

```
dat <- data.frame(variable,
                  coefficient = pls_coef,
                  comp)
```

```
ggplot(dat, aes(x = comp, y = coefficient, col = variable)) +
  geom_step() +
  xlab("# Component") +
  ylab("Coefficient") +
  theme(legend.position = "none")
```



4.3) Ridge Regression

```
ridgeGrid <- data.frame(.lambda = seq(0, .1, length = 15))
set.seed(1991)
ridge_fit <- train(x = solTrainXtrans, y = solTrainY,
  method = "ridge",
  tuneGrid = ridgeGrid,
  trControl = ctrl,
  preProcess = c("center", "scale"))
```


```
ridge_fit
```

```
## Ridge Regression
##
## 951 samples
## 228 predictors
##
## Pre-processing: centered (228), scaled (228)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 855, 857, 855, 857, 857, 855, ...
## Resampling results across tuning parameters:
##
##   lambda      RMSE      Rsquared    MAE
## 0.000000000 0.7071944 0.8814257 0.5257221
## 0.007142857 0.6919235 0.8869550 0.5236807
```



```
## 0.014285714 0.6856703 0.8890874 0.5203141
## 0.021428571 0.6836460 0.8898722 0.5199234
## 0.028571429 0.6835001 0.8900898 0.5213092
## 0.035714286 0.6844180 0.8900088 0.5230327
## 0.042857143 0.6860243 0.8897518 0.5248594
## 0.050000000 0.6881167 0.8893846 0.5270827
## 0.057142857 0.6905749 0.8889460 0.5295143
## 0.064285714 0.6933224 0.8884605 0.5321482
## 0.071428571 0.6963083 0.8879441 0.5348864
## 0.078571429 0.6994975 0.8874080 0.5378075
## 0.085714286 0.7028648 0.8868596 0.5407873
## 0.092857143 0.7063920 0.8863046 0.5438502
## 0.100000000 0.7100654 0.8857467 0.5469702
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was lambda = 0.02857143.
```

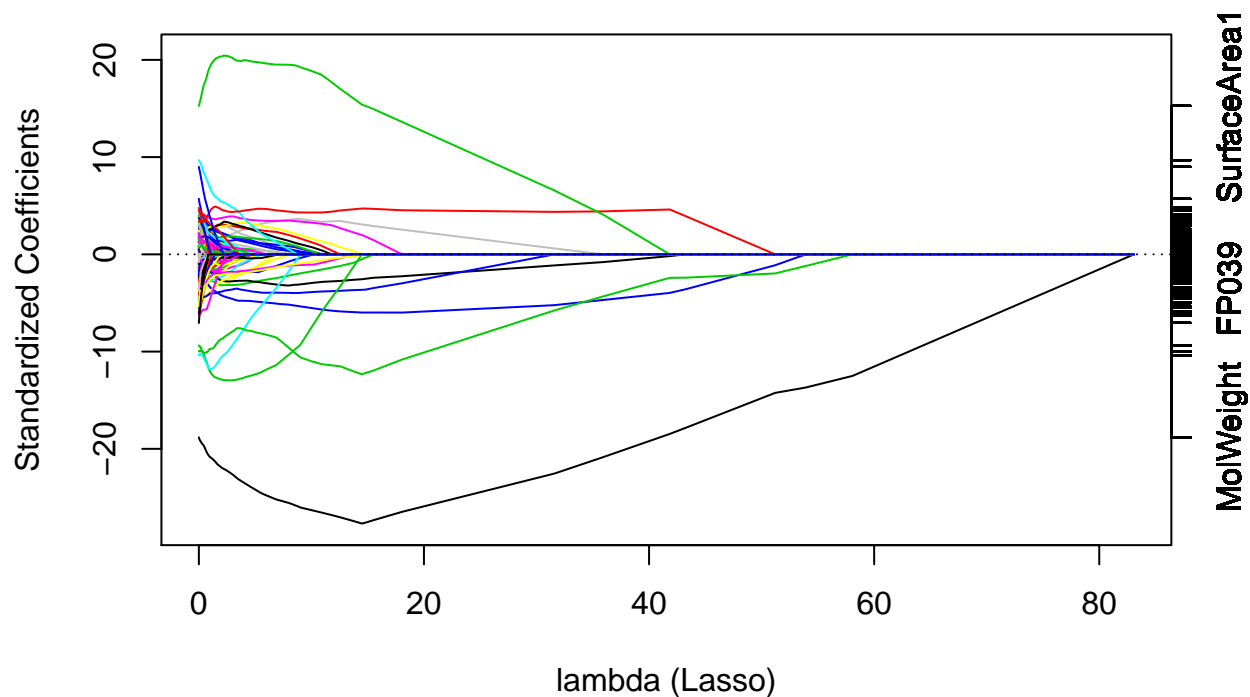
```
# plot the RMSEs against the values of the tuning parameter lambda
plot(ridge_fit)
```



```
# the value of lambda that gives the minimum RMSE value
ridge_fit$bestTune
```

```
##      lambda
## 5 0.02857143
```

```
# make a plot of the regression coefficient paths
plot(ridge_fit$finalModel, xvar = "penalty", use.color = TRUE)
```



4.4) Lasso

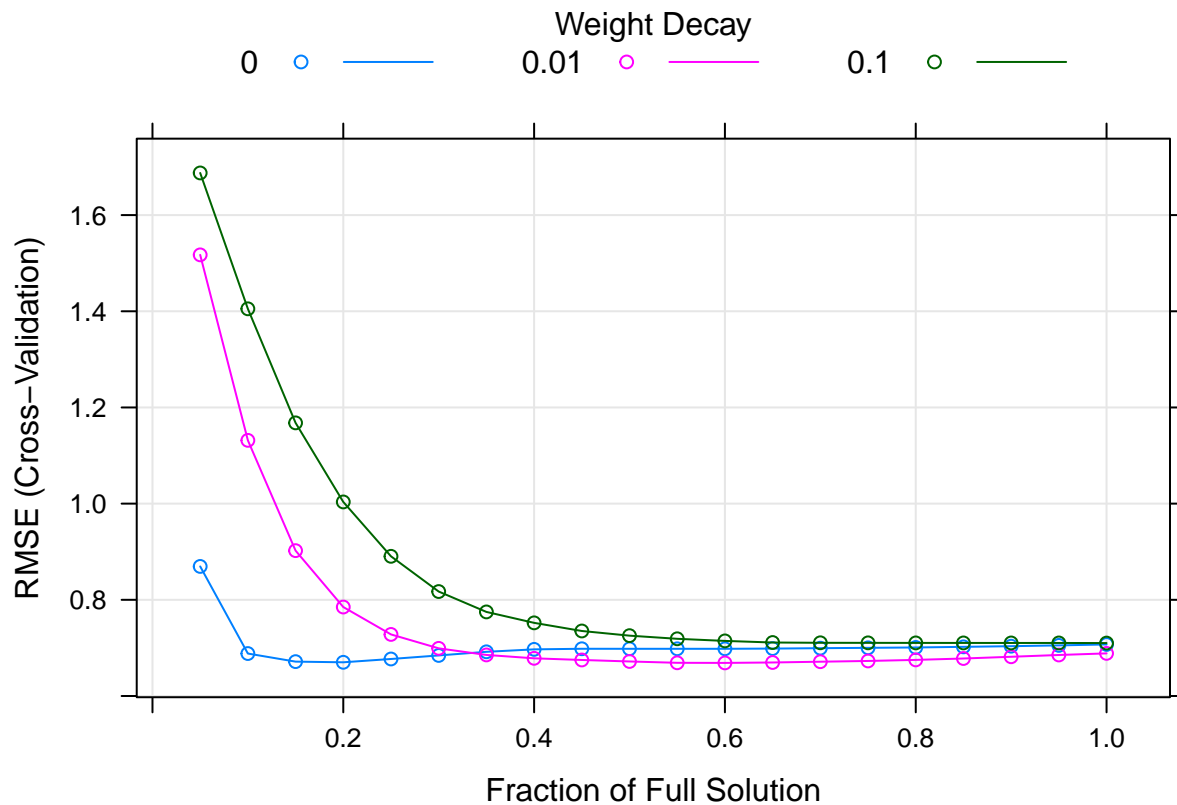
```
enetGrid <- expand.grid(.lambda = c(0, 0.01, .1),
                      .fraction = seq(.05, 1, length = 20))
set.seed(1991)
lasso_fit <- train(x = solTrainXtrans, y = solTrainY,
                  method = "enet",
                  tuneGrid = enetGrid,
                  trControl = ctrl,
                  preProcess = c("center", "scale"))
```

lasso_fit

```
## Elasticnet
##
## 951 samples
## 228 predictors
##
## Pre-processing: centered (228), scaled (228)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 855, 857, 855, 857, 857, 855, ...
## Resampling results across tuning parameters:
##
##  lambda  fraction  RMSE      Rsquared  MAE
##  0.00    0.05     0.8694393  0.8367154  0.6593669
##  0.00    0.10     0.6882519  0.8883257  0.5249954
##  0.00    0.15     0.6714866  0.8931908  0.5133845
##  0.00    0.20     0.6701005  0.8935851  0.5125415
##  0.00    0.25     0.6769459  0.8914827  0.5126865
##  0.00    0.30     0.6843558  0.8890745  0.5179973
```

##	0.00	0.35	0.6918879	0.8866338	0.5234601
##	0.00	0.40	0.6970105	0.8849352	0.5280080
##	0.00	0.45	0.6981910	0.8845023	0.5288807
##	0.00	0.50	0.6981295	0.8844524	0.5284216
##	0.00	0.55	0.6981444	0.8844082	0.5278100
##	0.00	0.60	0.6981516	0.8843850	0.5269660
##	0.00	0.65	0.6985908	0.8842072	0.5259632
##	0.00	0.70	0.6994292	0.8838955	0.5254815
##	0.00	0.75	0.7001222	0.8836637	0.5251417
##	0.00	0.80	0.7009164	0.8834005	0.5248366
##	0.00	0.85	0.7021745	0.8830045	0.5247800
##	0.00	0.90	0.7035777	0.8825485	0.5248372
##	0.00	0.95	0.7052453	0.8820344	0.5251640
##	0.00	1.00	0.7071944	0.8814257	0.5257221
##	0.01	0.05	1.5172575	0.6481926	1.1629621
##	0.01	0.10	1.1316121	0.7707836	0.8644227
##	0.01	0.15	0.9022421	0.8266589	0.6836986
##	0.01	0.20	0.7849595	0.8580220	0.5983956
##	0.01	0.25	0.7280635	0.8755480	0.5555230
##	0.01	0.30	0.6990802	0.8845478	0.5341814
##	0.01	0.35	0.6854148	0.8887916	0.5246724
##	0.01	0.40	0.6783782	0.8909363	0.5192493
##	0.01	0.45	0.6749279	0.8919645	0.5173129
##	0.01	0.50	0.6717883	0.8929277	0.5154451
##	0.01	0.55	0.6690558	0.8938298	0.5135199
##	0.01	0.60	0.6687896	0.8939771	0.5131305
##	0.01	0.65	0.6697231	0.8937514	0.5126997
##	0.01	0.70	0.6712940	0.8933490	0.5123954
##	0.01	0.75	0.6728448	0.8928764	0.5123212
##	0.01	0.80	0.6750877	0.8921643	0.5130300
##	0.01	0.85	0.6779499	0.8912939	0.5144052
##	0.01	0.90	0.6815677	0.8901898	0.5167259
##	0.01	0.95	0.6852845	0.8890690	0.5194156
##	0.01	1.00	0.6886007	0.8880743	0.5217469
##	0.10	0.05	1.6876421	0.5172217	1.2938453
##	0.10	0.10	1.4050887	0.7004963	1.0731807
##	0.10	0.15	1.1679369	0.7633412	0.8905843
##	0.10	0.20	1.0035064	0.7912991	0.7622082
##	0.10	0.25	0.8904595	0.8233943	0.6744600
##	0.10	0.30	0.8172256	0.8436821	0.6226784
##	0.10	0.35	0.7748367	0.8565541	0.5963412
##	0.10	0.40	0.7521196	0.8648159	0.5780088
##	0.10	0.45	0.7352072	0.8716584	0.5647150
##	0.10	0.50	0.7254667	0.8759264	0.5577237
##	0.10	0.55	0.7189654	0.8789894	0.5534675
##	0.10	0.60	0.7145841	0.8810958	0.5507771
##	0.10	0.65	0.7112406	0.8827524	0.5481277
##	0.10	0.70	0.7106350	0.8834483	0.5473125
##	0.10	0.75	0.7105176	0.8838941	0.5467777
##	0.10	0.80	0.7103955	0.8843254	0.5466359
##	0.10	0.85	0.7103572	0.8847166	0.5469637
##	0.10	0.90	0.7103120	0.8850839	0.5471206
##	0.10	0.95	0.7102830	0.8854133	0.5471128
##	0.10	1.00	0.7100654	0.8857467	0.5469702


```
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were fraction = 0.6 and lambda = 0.01.
# plot the RMSEs against the values of the tuning parameter lambda
plot(lasso_fit)
```



```
# the value of lambda that gives the minimum RMSE value
lasso_fit$bestTune

##      fraction lambda
## 32      0.6    0.01

# make a plot of the regression coefficient paths
plot(lasso_fit$finalModel, xvar = "penalty", use.color = TRUE)
```



4.5) Model Selection

```
# compute test-MSEs for PCR
y_pcr <- predict(pcr_fit, solTestXtrans)
MSE_pcr <- mean((solTestY - y_pcr)^2)
MSE_pcr
```

```
## [1] 0.8070032
```

```
# compute test-MSEs for PLSR
y_pls <- predict(pls_fit, solTestXtrans)
MSE_pls <- mean((solTestY - y_pls)^2)
MSE_pls
```

```
## [1] 0.5326949
```

```
# compute test-MSEs for Ridge Regression
y_ridge <- predict(ridge_fit, solTestXtrans)
MSE_ridge <- mean((solTestY - y_ridge)^2)
MSE_ridge
```

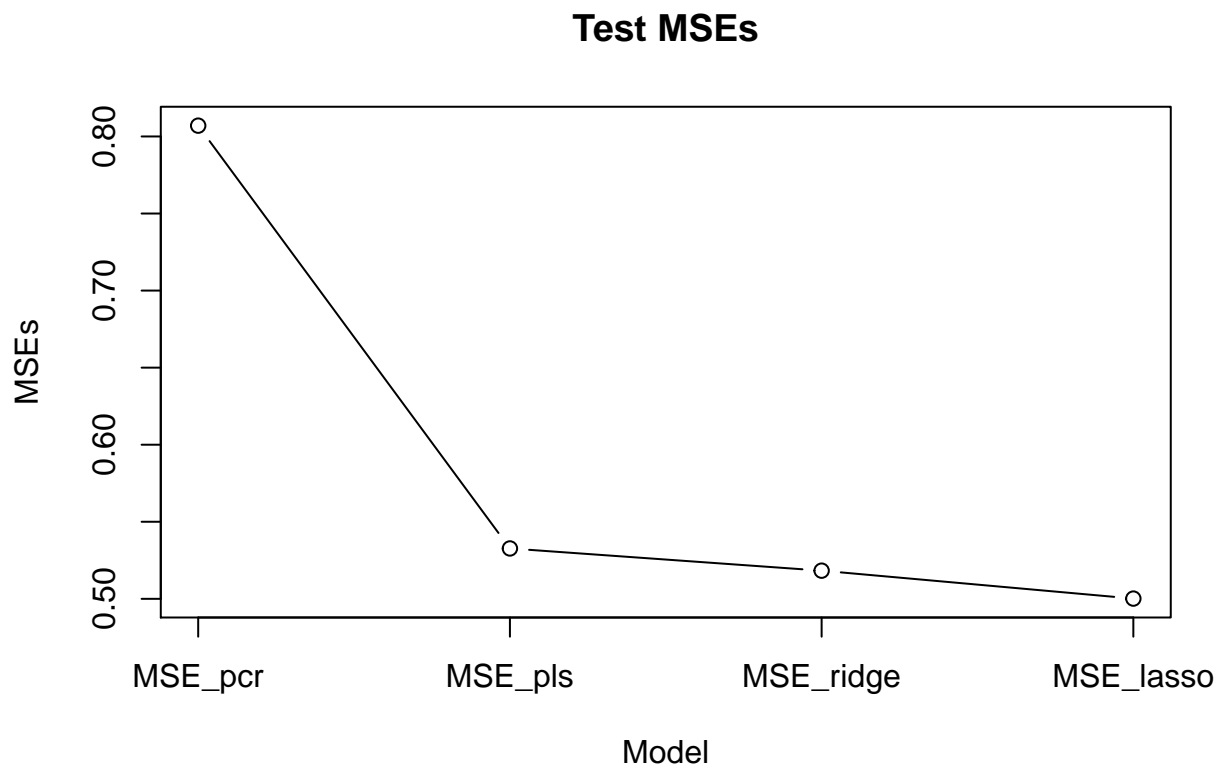
```
## [1] 0.5182821
```

```
# compute test-MSEs for PCR
y_lasso <- predict(lasso_fit, solTestXtrans)
MSE_lasso <- mean((solTestY - y_lasso)^2)
MSE_lasso
```

```
## [1] 0.5001638
```

```
# graph the test-MSEs
MSEs <- rbind(MSE_pcr, MSE_pls, MSE_ridge, MSE_lasso)
plot(MSEs, main = "Test MSEs", xlab = "Model", type = "b", xaxt = "n")
```

```
axis(1, at = 1:4, labels = rownames(MSEs))
```



```
# the smallest test-MSE  
MSEs[which.min(MSEs), 1]
```

```
## MSE_lasso  
## 0.5001638
```